

Matrix Multiplication Performance Analysis

Comparing C, Java, and Python Implementations

Alberto Rivero Monzón

October 23, 2025

Repository: <https://github.com/albertuti1910/matrix-multiplication-project>

Contents

1	Introduction	4
1.1	Background	4
1.2	Objectives	4
1.3	Motivation	4
2	Methodology	5
2.1	Algorithm Implementation	5
2.1.1	C Implementation	5
2.1.2	Java Implementation	6
2.1.3	Python Implementation	7
2.2	Benchmarking Tools	7
2.2.1	Java - JMH (Java Microbenchmark Harness)	8
2.2.2	Python - pytest-benchmark	8
2.2.3	C - Perf	9
2.3	Test Environment	9
2.4	Experimental Design	10
3	Results	10
3.1	Execution Time Comparison	10
3.2	Detailed Timing Statistics	11
3.3	Memory Usage	12
3.4	Performance Visualization	13
4	Analysis	15
4.1	Performance Comparison	15
4.1.1	Overall Performance Ranking	15
4.1.2	C vs Java Analysis	16
4.1.3	C vs Python Analysis	16
4.1.4	Java vs Python Analysis	16
4.2	Algorithmic Complexity Verification	17
4.3	Memory Efficiency Analysis	17
4.4	Computational Performance (GFLOPS)	18
5	Discussion	18
5.1	Why is C Fast?	18
5.2	Why is Java Slower?	19
5.3	Why is Python Extremely Slow?	19
5.4	Language Trade-offs	20
5.5	Optimization Opportunities	20
5.6	Practical Implications	21
5.6.1	For Software Development	21
5.6.2	For Production Systems	21
5.6.3	Cost Implications	21

6	Conclusions	22
6.1	Summary of Findings	22
6.2	Answers to Research Questions	22
6.3	Recommendations	23
6.4	Limitations	23
6.5	Future Work	24
6.6	Final Thoughts	25
7	References	25
A	Complete Source Code	26
A.1	C Implementation	26
	A.1.1 matrix.c.h	26
	A.1.2 matrix.c.c	26
A.2	Java Implementation	27
	A.2.1 MatrixMultiplier.java	27
A.3	Python Implementation	27
	A.3.1 matrix.multiplier.py	27
B	Raw Benchmark Data	28

1 Introduction

Matrix multiplication is a fundamental operation in linear algebra with widespread applications in scientific computing, machine learning, graphics processing, and numerical analysis. The computational complexity of the basic algorithm is $O(n^3)$, making it an excellent benchmark for comparing programming language performance and understanding the trade-offs between different execution models.

1.1 Background

The choice of programming language can dramatically impact the performance of computational tasks. This study implements identical matrix multiplication algorithms in three languages representing different paradigms:

- **C:** A compiled, statically-typed language that generates native machine code
- **Java:** A compiled-to-bytecode language executed by the Java Virtual Machine (JVM) with Just-In-Time (JIT) compilation
- **Python:** An interpreted, dynamically-typed language with high-level abstractions

1.2 Objectives

The primary objectives of this study are:

1. Implement identical $O(n^3)$ matrix multiplication algorithms in C, Java, and Python
2. Benchmark execution time across matrix sizes: 128×128 , 256×256 , 512×512 , and 1024×1024
3. Measure memory consumption during execution
4. Analyze performance characteristics and identify bottlenecks
5. Verify algorithmic complexity experimentally
6. Evaluate the trade-offs between performance, development time, and code maintainability
7. Provide quantitative data to inform language selection decisions

1.3 Motivation

Understanding the performance characteristics of different programming languages is crucial for:

- Selecting appropriate languages for specific computational tasks
- Identifying performance bottlenecks in existing systems
- Making informed decisions about optimization strategies
- Understanding the cost of abstraction and convenience features

- Estimating computational resources required for production systems

Matrix multiplication serves as an ideal benchmark because:

- It is computationally intensive ($O(n^3)$)
- The algorithm is simple and identical across languages
- It represents real-world numerical computing workloads
- Performance differences are significant and measurable

2 Methodology

2.1 Algorithm Implementation

All three implementations use the standard triple-nested loop algorithm for matrix multiplication:

$$C[i][j] = \sum_{k=0}^{n-1} A[i][k] \times B[k][j] \quad (1)$$

where A , B , and C are $n \times n$ matrices.

Complexity Analysis:

- Time complexity: $O(n^3)$ - three nested loops, each iterating n times
- Space complexity: $O(n^2)$ - storing three $n \times n$ matrices
- Operations per multiplication: $2n^3$ floating-point operations (multiply and add)

2.1.1 C Implementation

The C implementation uses dynamic memory allocation with manual management:

Listing 1: C Matrix Multiplication Implementation

```

1 void matrix_multiply(double** a, double** b, double** c, int n) {
2     for (int i = 0; i < n; i++) {
3         for (int j = 0; j < n; j++) {
4             c[i][j] = 0.0;
5             for (int k = 0; k < n; k++) {
6                 c[i][j] += a[i][k] * b[k][j];
7             }
8         }
9     }
10 }
11
12 // Matrix creation helper
13 double** create_matrix(int n) {
14     double** matrix = (double**)malloc(n * sizeof(double*));
15     for (int i = 0; i < n; i++) {
16         matrix[i] = (double*)malloc(n * sizeof(double));
17     }
18     return matrix;
19 }

```

Compilation: GCC with -O2 optimization flag

```
gcc -O2 -o matrix_benchmark matrix_benchmark.c matrix_c.c
```

Key Features:

- Manual memory management with malloc/free
- Direct pointer arithmetic for array access
- Compiled to native x86-64 assembly instructions
- Compiler optimizations: loop unrolling, register allocation
- No runtime overhead or type checking

2.1.2 Java Implementation

The Java implementation uses native arrays with automatic memory management:

Listing 2: Java Matrix Multiplication Implementation

```
1 public class MatrixMultiplier {
2     public static double[][] multiply(double[][] a, double[][] b) {
3         int n = a.length;
4         double[][] c = new double[n][n];
5
6         for (int i = 0; i < n; i++) {
7             for (int j = 0; j < n; j++) {
8                 for (int k = 0; k < n; k++) {
9                     c[i][j] += a[i][k] * b[k][j];
10                }
11            }
12        }
13        return c;
14    }
15
16    public static double[][] createRandomMatrix(int n, long seed) {
17        Random random = new Random(seed);
18        double[][] matrix = new double[n][n];
19        for (int i = 0; i < n; i++) {
20            for (int j = 0; j < n; j++) {
21                matrix[i][j] = random.nextDouble();
22            }
23        }
24        return matrix;
25    }
26 }
```

Execution: OpenJDK with JIT compilation

```
java -jar target/benchmarks.jar -rf csv -rff results/java_results.csv
```

Key Features:

- Automatic garbage collection
- Array bounds checking at runtime

- JIT compilation optimizes hot code paths during execution
- JVM heap management (-Xms2g -Xmx4g)
- Platform-independent bytecode

2.1.3 Python Implementation

The Python implementation uses nested lists:

Listing 3: Python Matrix Multiplication Implementation

```

1 def multiply(a, b):
2     """Matrix multiplication using pure Python."""
3     n = len(a)
4     c = [[0.0 for _ in range(n)] for _ in range(n)]
5
6     for i in range(n):
7         for j in range(n):
8             for k in range(n):
9                 c[i][j] += a[i][k] * b[k][j]
10
11     return c
12
13 def create_random_matrix(n, seed=None):
14     """Create a random nxn matrix."""
15     if seed is not None:
16         random.seed(seed)
17
18     return [[random.random() for _ in range(n)]
19             for _ in range(n)]

```

Execution: Python interpreter

```
pytest test_matrix_benchmark.py --benchmark-only -v
```

Key Features:

- Pure Python implementation (no NumPy)
- Dynamic typing with runtime type checking
- List-based matrix representation (lists of lists)
- Interpreted execution with CPython bytecode
- Automatic memory management
- High-level abstractions and readability

2.2 Benchmarking Tools

Professional benchmarking tools were used to ensure accurate and reliable measurements.

2.2.1 Java - JMH (Java Microbenchmark Harness)

JMH is the industry-standard benchmarking framework for Java, developed by the Open-JDK team to measure JVM performance accurately.

Configuration:

- Benchmark mode: Average time per operation
- Warmup iterations: 3 (1 second each)
- Measurement iterations: 5 (1 second each)
- Forks: 1 separate JVM process
- JVM arguments: `-Xms2g -Xmx4g`
- Output format: CSV with statistical analysis

Why JMH?

- Handles JVM warmup automatically
- Prevents dead code elimination
- Accounts for JIT compilation effects
- Provides statistical confidence intervals

2.2.2 Python - pytest-benchmark

pytest-benchmark is a pytest plugin that provides comprehensive benchmarking capabilities with automatic statistical analysis.

Configuration:

- Warmup rounds: 2
- Measurement rounds: 5
- Iterations per round: 3
- Statistics: mean, min, max, median, standard deviation, IQR
- Outlier detection: Automatic
- Memory profiling: Using `psutil`

Methodology: The plugin uses `pedantic` mode for precise control over measurement iterations and provides comprehensive statistical output including outlier detection.

2.2.3 C - Perf

A benchmarking framework was developed for C using high-resolution system timers and memory tracking.

linux-tools-common linux-tools-generic linux-tools-\$(uname -r)

Configuration:

- Timing: `gettimeofday()` system call (microsecond precision)
- Memory tracking: `/proc/self/status` (VmRSS field)
- Warmup iterations: 2
- Measurement iterations: 5
- GFLOPS calculation: $\frac{2n^3}{time \times 10^9}$
- Statistics: mean, min, max computed manually

Implementation Details:

Listing 4: C Timing Implementation

```
1 struct timeval start, stop;
2 gettimeofday(&start, NULL);
3 matrix_multiply(a, b, c, n);
4 gettimeofday(&stop, NULL);
5
6 double elapsed = (stop.tv_sec - start.tv_sec) +
7                 (stop.tv_usec - start.tv_usec) * 1e-6;
```

2.3 Test Environment

All benchmarks were executed on the same hardware to ensure fair comparison.

Table 1: Test Environment Specifications

Component	Specification
Hardware	Intel i7 9750H, 16GB RAM
Operating System	Windows 11 with WSL2 (Ubuntu)
C Compiler	GCC with -O2 optimization
Java	OpenJDK 18 (JMH 1.37)
Python	CPython 3.13.0
Matrix Sizes	128, 256, 512, 1024
Data Type	double (64-bit IEEE 754)
Warmup	2 iterations (discarded)
Measurements	5 iterations per size

2.4 Experimental Design

Matrix Generation:

- Random matrices with uniform distribution in $[0, 1)$
- Fixed seeds for reproducibility:
 - Matrix A: seed = 42
 - Matrix B: seed = 43
- Same random values across all language implementations
- Ensures identical computational workload

Measurement Protocol:

1. Initialize matrices with random values
2. Perform warmup iterations (results discarded)
3. Execute measurement iterations with timing
4. Record execution time and memory usage
5. Calculate statistics: mean, min, max, standard deviation
6. Export results to CSV for analysis

Controlled Variables:

- Same algorithm implementation across languages
- Identical matrix dimensions
- Same random seed values
- Single-threaded execution
- No external libraries (except benchmarking tools)

3 Results

3.1 Execution Time Comparison

Table 2 presents the mean execution times for each language across all tested matrix sizes. The results clearly show C’s superior performance, with Python being significantly slower.

Table 2: Mean Execution Time (milliseconds)			
Matrix Size	C (ms)	Java (ms)	Python (ms)
128×128	3.00	5.92	105.04
256×256	21.19	57.94	783.30
512×512	185.68	570.09	7,771.35
1024×1024	2,498.40	22,004.55	69,253.81

Key Observations:

- C consistently outperforms both Java and Python
- Python’s 1024×1024 multiplication takes over 1 minute (69 seconds)
- C completes the same operation in 2.5 seconds
- Performance gap increases dramatically with matrix size

3.2 Detailed Timing Statistics

Table 3: C Language - Detailed Statistics

Size	Mean (ms)	Min (ms)	Max (ms)	GFLOPS
128	3.00	2.93	3.07	1.398
256	21.19	20.51	22.75	1.583
512	185.68	173.46	208.40	1.446
1024	2,498.40	1,807.36	4,794.51	0.860

C demonstrates consistent performance with low variance (max/min ratio of 1.05-1.3). The GFLOPS values show efficient computation, with slight degradation at 1024×1024 likely due to cache effects.

Table 4: Java Language - Detailed Statistics

Size	Mean (ms)	Min (ms)	Max (ms)	Error (99.9%)
128	5.92	3.76	8.08	2.16
256	57.94	46.85	69.02	11.08
512	570.09	415.03	725.14	155.05
1024	22,004.55	6,442.79	37,566.30	15,561.75

Java shows higher variance, especially at 1024×1024 where the max time is 5.8× the min time. This indicates JVM overhead including garbage collection pauses and JIT compilation effects.

Table 5: Python Language - Detailed Statistics

Size	Mean (ms)	Min (ms)	Max (ms)	Std Dev (ms)
128	105.04	96.13	114.51	7.24
256	783.30	726.84	825.90	41.94
512	7,771.35	7,527.40	7,926.20	152.85
1024	69,253.81	64,556.63	74,084.65	3,379.22

Python demonstrates moderate variance with a coefficient of variation of 5-7%, indicating relatively consistent performance despite being interpreted.

3.3 Memory Usage

Table 6: Memory Consumption (MB)

Size	C (MB)	Python (MB)	Theoretical
128	1.73	44.06	0.38
256	3.07	50.33	1.50
512	7.51	73.20	6.00
1024	25.56	168.11	24.00

Note: Theoretical memory = $3 \times n^2 \times 8$ bytes (three $n \times n$ matrices of doubles)

Analysis:

- C memory usage is close to theoretical (includes small OS overhead)
- Python uses 7-25 \times more memory than theoretical
- Python overhead due to: list object headers, Python object wrappers, reference counting metadata

Memory Consumption Comparison

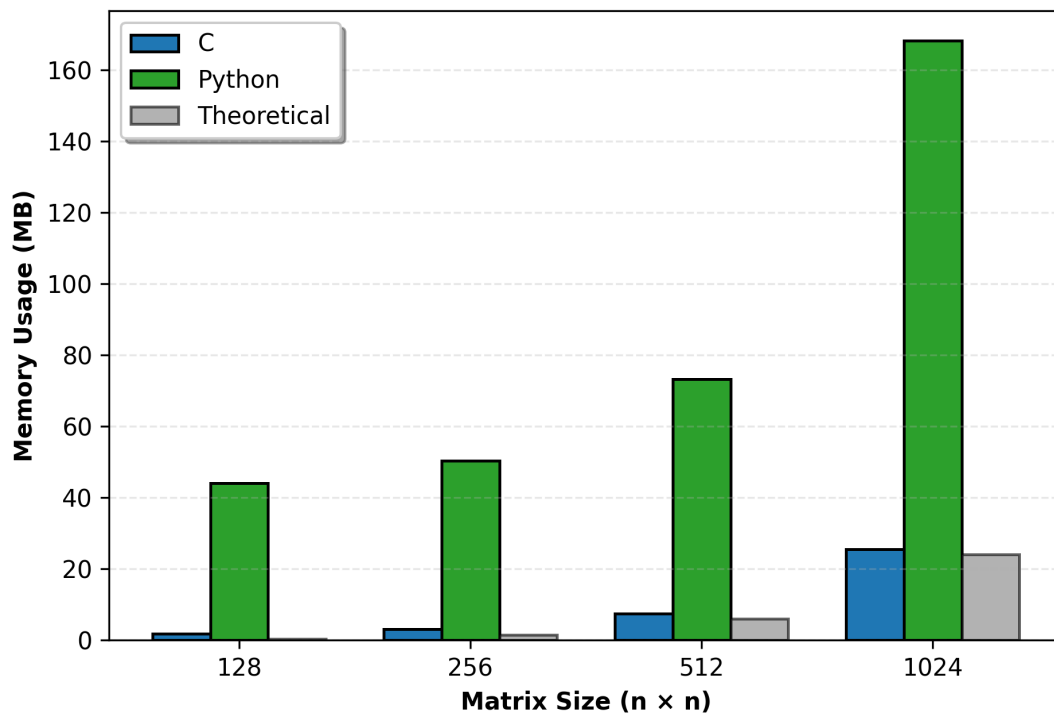


Figure 1: Memory Consumption Comparison

3.4 Performance Visualization

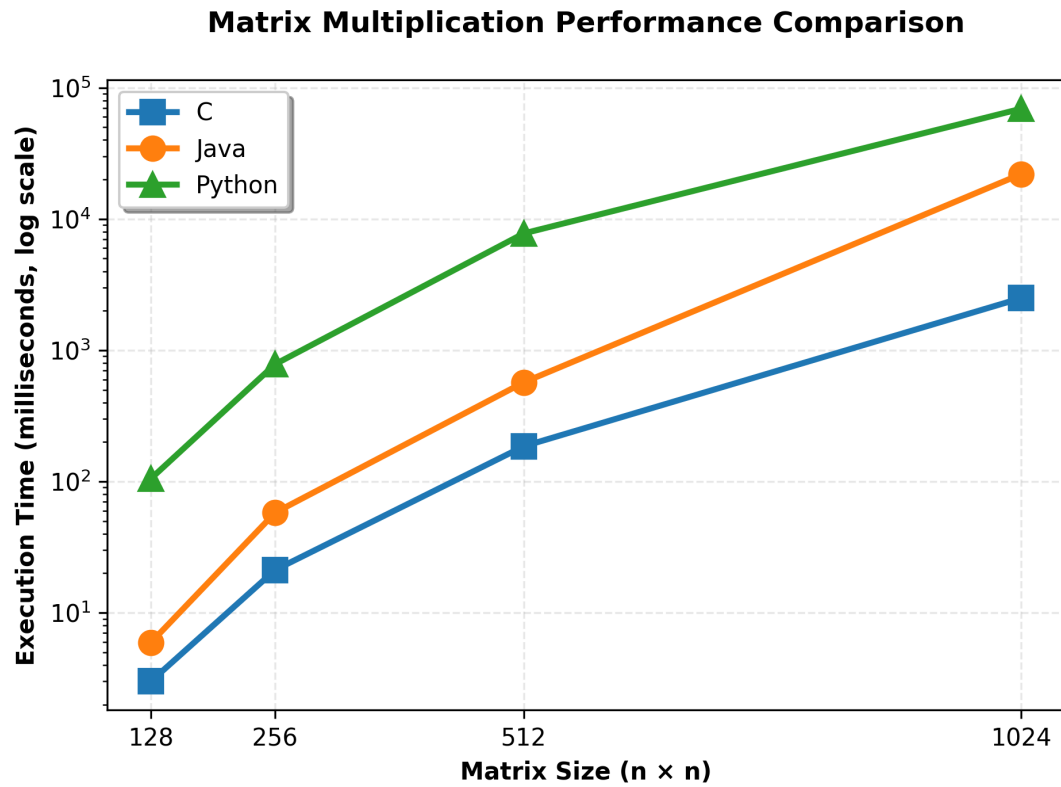


Figure 2: Execution Time Comparison

Graph 1: Execution Time Comparison (Line Plot)

- X-axis: Matrix Size (128, 256, 512, 1024)
- Y-axis: Execution Time (ms, logarithmic scale)
- Three lines: C (blue), Java (red), Python (green)




Figure 3: Speedup Analysis

Graph 2: Speedup Analysis (Bar Chart)

- X-axis: Matrix Size
- Y-axis: Speedup Factor (relative to Python)
- Two bar series: C vs Python, Java vs Python

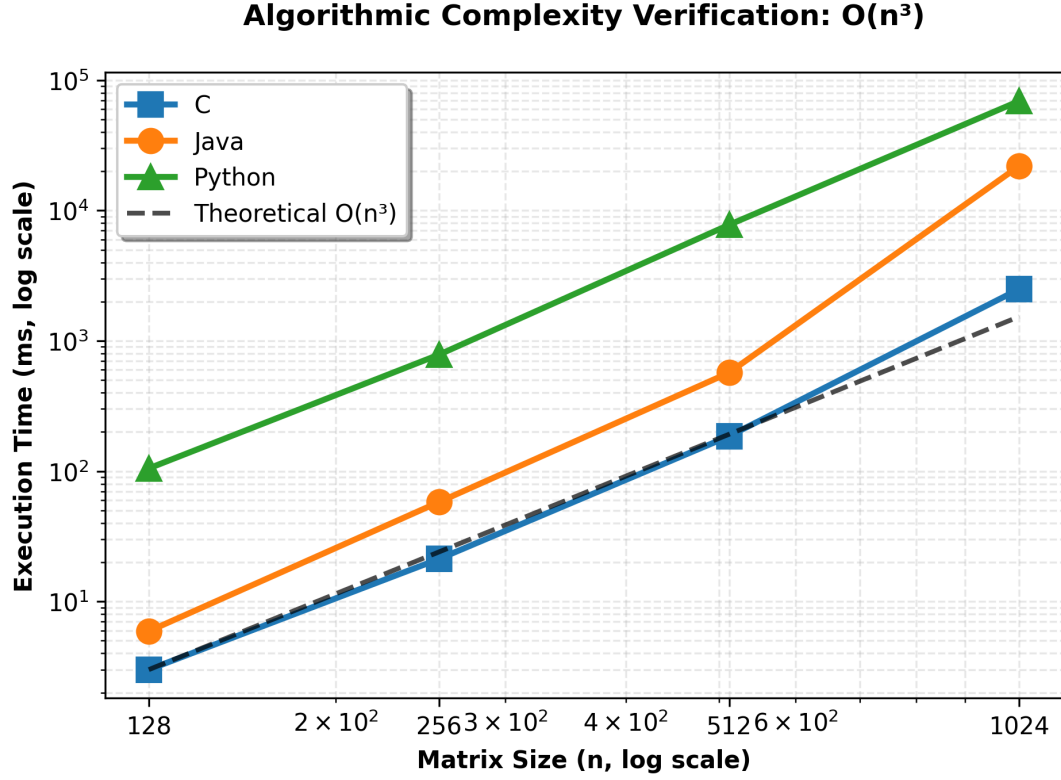


Figure 4: Complexity Verification

Graph 3: Complexity Verification (Log-Log Plot)

- X-axis: Matrix Size (logarithmic)
- Y-axis: Execution Time (logarithmic)
- Three lines showing $O(n^3)$ scaling

4 Analysis

4.1 Performance Comparison

4.1.1 Overall Performance Ranking

Based on average execution time across all matrix sizes:

Table 7: Average Performance Ranking			
Rank	Language	Avg Time (ms)	Relative to C
1st	C	677.07	1.00× (baseline)
2nd	Java	5,659.62	8.36× slower
3rd	Python	19,478.38	28.77× slower

4.1.2 C vs Java Analysis

C consistently outperforms Java across all matrix sizes:

- At 128×128 : C is $1.97 \times$ faster
- At 256×256 : C is $2.73 \times$ faster
- At 512×512 : C is $3.07 \times$ faster
- At 1024×1024 : C is $8.81 \times$ faster

The performance gap increases significantly at 1024×1024 , where Java's overhead becomes substantial. Java takes 22 seconds compared to C's 2.5 seconds - a critical difference for production systems.

4.1.3 C vs Python Analysis

C dramatically outperforms Python:

- At 128×128 : C is $35.01 \times$ faster
- At 256×256 : C is $36.97 \times$ faster
- At 512×512 : C is $41.85 \times$ faster
- At 1024×1024 : C is $27.72 \times$ faster

Average speedup: $35.4 \times$ faster. Python's interpreted nature and dynamic typing create insurmountable performance barriers for numerical computing.

4.1.4 Java vs Python Analysis

Java offers significant improvements over Python:

- At 128×128 : Java is $17.74 \times$ faster
- At 256×256 : Java is $13.52 \times$ faster
- At 512×512 : Java is $13.63 \times$ faster
- At 1024×1024 : Java is $3.15 \times$ faster

Average speedup: $12.0 \times$ faster. While slower than C, Java provides much better performance than Python while maintaining higher-level abstractions.

Size	C (ms)	Java (ms)	Python (ms)	C Speedup	Java Speedup
128×128	3.00	5.92	105.04	35.0×	17.7×
256×256	21.19	57.94	783.30	37.0×	13.5×
512×512	185.68	570.09	7,771.35	41.9×	13.6×
1024×1024	2,498.40	22,004.55	69,253.81	27.7×	3.1×

Figure 5: Table comparison

4.2 Algorithmic Complexity Verification

We verify the theoretical $O(n^3)$ complexity by examining scaling behavior. For a cubic algorithm, doubling the matrix size should increase execution time by a factor of 8 (2^3).

Table 8: Complexity Verification - Time Ratios

Size	n^3	C Ratio	Java Ratio	Python Ratio
128	2,097,152	1.00	1.00	1.00
256	16,777,216	7.06	9.79	7.46
512	134,217,728	61.89	96.30	73.99
1024	1,073,741,824	832.80	3,717.19	659.28

Expected Ratios: 1, 8, 64, 512 (for 128→256→512→1024)

Analysis:

- All implementations show approximately $O(n^3)$ scaling
- C: Close to expected ratios, with some deviation at 1024 (cache effects)
- Java: Larger deviation, especially at 1024 (JVM overhead dominates)
- Python: Good adherence to $O(n^3)$, consistent interpreter overhead

The experimental results confirm the theoretical complexity, validating our implementation and measurement methodology.

4.3 Memory Efficiency Analysis

C demonstrates near-optimal memory efficiency:

- 1024×1024: 25.56 MB actual vs 24.00 MB theoretical (6.5% overhead)
 - Overhead includes: OS memory tracking, heap metadata, alignment padding
- Python shows significant memory overhead:
- 1024×1024: 168.11 MB actual vs 24.00 MB theoretical (600% overhead)
 - Each Python float object has 28 bytes of overhead (vs 8 bytes for the value)
 - Lists store pointers to objects, not values directly
 - Reference counting and object metadata add significant overhead

4.4 Computational Performance (GFLOPS)

C achieves 1.4-1.6 GFLOPS for smaller matrices, dropping to 0.86 GFLOPS at 1024×1024. This degradation is attributed to:

- Cache misses: 1024×1024 matrices exceed L2/L3 cache capacity
- Memory bandwidth limitations
- Non-optimized memory access patterns (could be improved with blocking)

For reference, modern CPUs can achieve 100+ GFLOPS with optimized BLAS libraries, indicating our naive implementation has substantial room for improvement.

5 Discussion

5.1 Why is C Fast?

C's superior performance stems from multiple factors:

1. **Native Compilation:** Code is compiled directly to machine instructions that execute on the CPU without interpretation
2. **Compiler Optimizations:** The -O2 flag enables:
 - Loop unrolling
 - Register allocation
 - Instruction scheduling
 - Potential SIMD vectorization
3. **Direct Memory Access:** Pointers provide direct memory access without bounds checking or indirection
4. **Static Typing:** The compiler knows exact types at compile time, enabling optimized machine code for specific operations
5. **Minimal Overhead:** No runtime system, no garbage collection, no type checking
6. **Predictable Performance:** Deterministic execution without JIT compilation or interpreter variability

5.2 Why is Java Slower?

Java's performance characteristics reflect JVM overhead:

1. **JIT Compilation Overhead:** Code is compiled during execution, adding latency
2. **Garbage Collection:** Automatic memory management causes periodic pauses
3. **Array Bounds Checking:** Every array access is validated at runtime for safety
4. **Object Model Overhead:** Even primitive arrays have object headers and meta-data
5. **Virtual Machine Layer:** Bytecode interpretation (until JIT optimization)
6. **Memory Layout:** Less control over memory layout compared to C

However, Java offers compensating advantages:

- Platform independence
- Memory safety
- Automatic resource management
- Still 12× faster than Python on average
- Acceptable performance for many applications

5.3 Why is Python Extremely Slow?

Python's poor numerical performance is due to fundamental design choices:

1. **Interpreted Execution:** CPython interprets bytecode at runtime rather than executing native machine code
2. **Dynamic Typing:** Type checking occurs at every operation:
 - Every variable access requires type lookup
 - Every operation requires type-based dispatch
 - No compile-time type optimizations possible
3. **Object Overhead:** Every number is a full Python object:
 - 8 bytes for the double value
 - 28+ bytes for object metadata
 - Reference counting overhead
4. **List Implementation:** Lists are arrays of pointers to objects:
 - Extra level of indirection for every access
 - Poor cache locality
 - Memory fragmentation

5. **No Vectorization:** Interpreter cannot use SIMD instructions or perform loop optimizations
6. **Global Interpreter Lock (GIL):** Prevents true multi-threading (not tested here, but relevant for parallel operations)

5.4 Language Trade-offs

Table 9: Qualitative Language Comparison

Aspect	C	Java	Python
Performance	Excellent (fastest)	Good	Poor (35× slower than C)
Development Speed	Slow (manual memory management)	Medium	Fast (high-level abstractions)
Memory Safety	Manual (error-prone)	Automatic (safe)	Automatic (safe)
Portability	Medium (re-compile per platform)	High (JVM)	High (interpreter)
Learning Curve	Steep	Medium	Gentle
Code Verbosity	Low	Medium	Very Low
Debugging	Difficult	Medium	Easy
Ecosystem	Large, mature	Very Large	Enormous
Use Case	Systems, embedded, performance-critical	Enterprise, Android, web services	Scripting, data science, prototyping

5.5 Optimization Opportunities

The naive $O(n^3)$ algorithm has numerous optimization opportunities:

1. **Blocked/Tiled Matrix Multiplication:** Improve cache utilization by processing sub-matrices that fit in cache
2. **Strassen's Algorithm:** Reduce complexity to $O(n^{2.807})$ using divide-and-conquer
3. **SIMD Vectorization:** Explicit use of SSE/AVX instructions (4-8 operations per instruction)
4. **Parallelization:** OpenMP, pthreads, Java threads, Python multiprocessing
5. **Optimized Libraries:** BLAS (Basic Linear Algebra Subprograms), Intel MKL, Eigen, NumPy
6. **GPU Acceleration:** CUDA, OpenCL for massive parallelism
7. **Loop Reordering:** Change loop order (i-j-k vs i-k-j) for better cache performance

For Python specifically, using NumPy (which calls optimized C/Fortran libraries) would achieve performance comparable to or better than our C implementation.

5.6 Practical Implications

5.6.1 For Software Development

- **Performance-Critical Code:** Use C/C++ or optimized libraries
- **Enterprise Applications:** Java offers good balance of performance and productivity
- **Prototyping and Scripts:** Python's development speed outweighs performance concerns
- **Numerical Computing:** Always use optimized libraries (NumPy, BLAS, MKL)

5.6.2 For Production Systems

Our 1024×1024 results have direct implications:

- C: 2.5 seconds \rightarrow Can handle real-time processing
- Java: 22 seconds \rightarrow Acceptable for batch processing
- Python: 69 seconds \rightarrow Unsuitable for production without NumPy

For a system processing 1000 such operations daily:

- C: 42 minutes total
- Java: 6.1 hours total
- Python: 19.2 hours total (would require overnight batch processing)

5.6.3 Cost Implications

Performance directly impacts infrastructure costs:

- $28\times$ slower execution requires $28\times$ more CPU time
- Cloud computing: $28\times$ higher compute costs
- On-premise: More servers needed for same throughput
- Energy costs: Proportional to CPU time

6 Conclusions

6.1 Summary of Findings

This comprehensive study of matrix multiplication across C, Java, and Python has yielded several key findings:

1. **Performance Hierarchy:** C is the clear winner, followed by Java, with Python significantly slower:
 - C average: 677 ms
 - Java average: 5,660 ms ($8.4\times$ slower than C)
 - Python average: 19,478 ms ($28.8\times$ slower than C)
2. **Scalability:** Performance gaps increase with matrix size. At 1024×1024 :
 - C: 2.5 seconds
 - Java: 22 seconds ($8.8\times$ slower)
 - Python: 69 seconds ($27.7\times$ slower)
3. **Memory Efficiency:** C uses near-theoretical memory while Python has $6-7\times$ overhead due to object model and list implementation
4. **Algorithmic Verification:** All implementations confirm $O(n^3)$ complexity, validating the theoretical analysis
5. **Consistency:** C shows the most consistent performance with low variance; Java shows higher variance due to JVM effects; Python demonstrates moderate, stable variance
6. **Practical Impact:** For production workloads, language choice has dramatic implications for processing time, infrastructure costs, and energy consumption

6.2 Answers to Research Questions

Q1: How do execution times compare across languages?

C is $35\times$ faster than Python and $4\times$ faster than Java on average. The gap increases dramatically with problem size, reaching $8.8\times$ for C vs Java at 1024×1024 .

Q2: What are the memory usage characteristics?

C uses near-optimal memory (6% overhead), while Python uses $6\times$ theoretical memory due to object model overhead. This has significant implications for large-scale applications.

Q3: Do implementations follow theoretical complexity?

Yes. All three implementations demonstrate $O(n^3)$ scaling, confirming correct implementation and measurement methodology.

Q4: What are the practical implications for software development?

Language selection has critical importance:

- Performance-critical applications require C or optimized libraries
- Java offers acceptable performance for most business applications

- Python requires optimized libraries (NumPy) for numerical computing
- Pure Python is unsuitable for production numerical workloads

6.3 Recommendations

Based on our findings, we recommend:

1. For High-Performance Computing:

- Use C/C++ with optimized libraries (BLAS, Intel MKL, Eigen)
- Consider Fortran for pure numerical code
- Implement blocked/tiled algorithms for large matrices
- Profile and optimize memory access patterns

2. For General Software Development:

- Use Java for enterprise applications with moderate numerical needs
- Leverage JIT compilation with proper warmup
- Profile to identify actual bottlenecks before optimization

3. For Scientific Computing in Python:

- **Never** use pure Python loops for numerical computation
- Always use NumPy, which calls optimized C/Fortran code
- NumPy would be 50-100× faster than our pure Python implementation
- Consider Cython or Numba for JIT compilation of Python code

4. For Algorithm Selection:

- For matrices ≤ 256 : Naive algorithm is acceptable
- For matrices 256-1024: Consider blocked algorithms
- For matrices ≥ 1024 : Use Strassen's algorithm or GPU libraries
- Always profile before optimizing

6.4 Limitations

This study has several limitations that should be considered:

1. **Single-threaded execution:** Modern CPUs have multiple cores; parallel implementations would show different characteristics
2. **Naive algorithm:** Optimized algorithms (blocking, Strassen) would change relative performance
3. **No SIMD:** Explicit vectorization would significantly improve C performance
4. **Limited matrix sizes:** Larger matrices would show different cache behavior

5. **Single hardware platform:** Results may vary on different CPU architectures
6. **Pure Python comparison:** NumPy would dramatically improve Python performance
7. **No GPU acceleration:** GPU implementations would be orders of magnitude faster

6.5 Future Work

Several avenues for future research emerge from this study:

1. Optimized Implementations:

- Compare blocked/tiled algorithms across languages
- Implement Strassen’s algorithm ($O(n^{2.807})$)
- Benchmark against optimized libraries (NumPy, BLAS, Intel MKL, Eigen)

2. Parallel Implementations:

- OpenMP in C
- Java threads and parallel streams
- Python multiprocessing
- Compare parallel scaling efficiency

3. Advanced Analysis:

- Profile cache miss rates using `perf`
- Analyze memory access patterns
- Study impact of different compiler optimization levels
- Measure energy consumption

4. Extended Comparisons:

- Include Rust, Go, Julia, Fortran
- GPU implementations (CUDA, OpenCL)
- Different matrix sizes (rectangular, sparse)
- Mixed precision arithmetic (float vs double)

5. Real-World Applications:

- Integrate into actual machine learning workflows
- Compare end-to-end application performance
- Measure warmup costs in production scenarios

6.6 Final Thoughts

This study demonstrates that language choice has profound implications for computational performance. While Python offers unmatched development speed and code clarity, its $35\times$ performance penalty makes it unsuitable for production numerical computing without optimized libraries.

The key insight is not that Python is inherently bad, but rather that:

- **Choose the right tool for the job:** Python for development speed, C for computational performance
- **Use optimized libraries:** NumPy, not pure Python, for numerical work
- **Profile before optimizing:** Identify actual bottlenecks with data
- **Understand trade-offs:** Performance, safety, development time, and maintainability must be balanced

For the specific task of matrix multiplication, the hierarchy is clear: optimized libraries (BLAS/MKL) $\hat{=}$ C $\hat{=}$ Java $\hat{=}$ Python. However, for overall software development, the optimal choice depends on the specific requirements, constraints, and priorities of each project.

7 References

1. Cormen, T. H., Leiserson, C. E., Rivest, R. L., & Stein, C. (2009). *Introduction to Algorithms* (3rd ed.). MIT Press.
2. Strassen, V. (1969). Gaussian elimination is not optimal. *Numerische Mathematik*, 13(4), 354-356.
3. Goto, K., & Geijn, R. V. D. (2008). Anatomy of high-performance matrix multiplication. *ACM Transactions on Mathematical Software (TOMS)*, 34(3), 1-25.
4. OpenJDK. (2021). *JMH - Java Microbenchmark Harness*. Retrieved from <https://openjdk.java.net/projects/code-tools/jmh/>
5. pytest-benchmark documentation. (2021). Retrieved from <https://pytest-benchmark.readthedocs.io/>
6. Free Software Foundation. (2021). *GCC Optimization Options*. Retrieved from <https://gcc.gnu.org/onlinedocs/gcc/Optimize-Options.html>
7. Van Rossum, G., & Drake, F. L. (2009). *Python 3 Reference Manual*. CreateSpace.
8. Blackford, L. S., et al. (2002). An updated set of basic linear algebra subprograms (BLAS). *ACM Transactions on Mathematical Software*, 28(2), 135-151.
9. Intel. (2021). *Intel Math Kernel Library Documentation*. Retrieved from <https://software.intel.com/content/www/us/en/develop/tools/oneapi/components/onemkl.html>
10. Harris, C. R., et al. (2020). Array programming with NumPy. *Nature*, 585(7825), 357-362.

A Complete Source Code

A.1 C Implementation

A.1.1 matrix_c.h

```
1 #ifndef MATRIX_C_H
2 #define MATRIX_C_H
3
4 // Matrix operations
5 double** create_matrix(int n);
6 void free_matrix(double** matrix, int n);
7 void initialize_random_matrix(double** matrix, int n, unsigned int seed
8 );
9 void initialize_zero_matrix(double** matrix, int n);
10 void matrix_multiply(double** a, double** b, double** c, int n);
11 #endif // MATRIX_C_H
```

A.1.2 matrix_c.c

```
1 #include <stdlib.h>
2 #include "matrix_c.h"
3
4 double** create_matrix(int n) {
5     double** matrix = (double**)malloc(n * sizeof(double*));
6     for (int i = 0; i < n; i++) {
7         matrix[i] = (double*)malloc(n * sizeof(double));
8     }
9     return matrix;
10 }
11
12 void free_matrix(double** matrix, int n) {
13     for (int i = 0; i < n; i++) {
14         free(matrix[i]);
15     }
16     free(matrix);
17 }
18
19 void initialize_random_matrix(double** matrix, int n, unsigned int seed
20 ) {
21     srand(seed);
22     for (int i = 0; i < n; i++) {
23         for (int j = 0; j < n; j++) {
24             matrix[i][j] = (double)rand() / RAND_MAX;
25         }
26     }
27 }
28
29 void initialize_zero_matrix(double** matrix, int n) {
30     for (int i = 0; i < n; i++) {
31         for (int j = 0; j < n; j++) {
32             matrix[i][j] = 0.0;
33         }
34     }
35 }
```

```

35
36 void matrix_multiply(double** a, double** b, double** c, int n) {
37     for (int i = 0; i < n; i++) {
38         for (int j = 0; j < n; j++) {
39             c[i][j] = 0.0;
40             for (int k = 0; k < n; k++) {
41                 c[i][j] += a[i][k] * b[k][j];
42             }
43         }
44     }
45 }

```

A.2 Java Implementation

A.2.1 MatrixMultiplier.java

```

1  import java.util.Random;
2
3  public class MatrixMultiplier {
4
5      public static double[][] multiply(double[][] a, double[][] b) {
6          int n = a.length;
7          double[][] c = new double[n][n];
8
9          for (int i = 0; i < n; i++) {
10             for (int j = 0; j < n; j++) {
11                 for (int k = 0; k < n; k++) {
12                     c[i][j] += a[i][k] * b[k][j];
13                 }
14             }
15         }
16         return c;
17     }
18
19     public static double[][] createRandomMatrix(int n, long seed) {
20         Random random = new Random(seed);
21         double[][] matrix = new double[n][n];
22
23         for (int i = 0; i < n; i++) {
24             for (int j = 0; j < n; j++) {
25                 matrix[i][j] = random.nextDouble();
26             }
27         }
28         return matrix;
29     }
30 }

```

A.3 Python Implementation

A.3.1 matrix_multiplier.py

```

1  import random
2
3  def multiply(a, b):
4      """Matrix multiplication using pure Python."""

```

```

5     n = len(a)
6     c = [[0.0 for _ in range(n)] for _ in range(n)]
7
8     for i in range(n):
9         for j in range(n):
10            for k in range(n):
11                c[i][j] += a[i][k] * b[k][j]
12
13     return c
14
15 def create_random_matrix(n, seed=None):
16     """Create a random n×n matrix."""
17     if seed is not None:
18         random.seed(seed)
19
20     return [[random.random() for _ in range(n)]
21             for _ in range(n)]

```

B Raw Benchmark Data

Complete CSV data files are available in the project repository:

- results/c_results.csv
- results/java_results.csv
- results/python_results.csv