

# **Parallel (and Vectorized) Matrix Multiplication**

## Comparative Study

Alberto Rivero Monzón

December 4, 2025

Repository: <https://github.com/albertuti1910/matrix-multiplication-project>

# Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
<b>2</b>	<b>Theoretical Background</b>	<b>3</b>
2.1	OpenMP and Multi-threading . . . . .	3
2.2	SIMD (Single Instruction, Multiple Data) . . . . .	3
2.3	Cache Locality and Transposition . . . . .	3
<b>3</b>	<b>Methodology</b>	<b>4</b>
<b>4</b>	<b>Quantitative Results</b>	<b>4</b>
4.1	Execution Time . . . . .	4
4.2	Computational Throughput (GFLOPS) . . . . .	4
<b>5</b>	<b>Performance Analysis</b>	<b>5</b>
5.1	Execution Time and Overhead . . . . .	5
5.2	Analysis of the "Parallel Trap" . . . . .	6
<b>6</b>	<b>Conclusion</b>	<b>7</b>

# 1 Introduction

In this assignment, we investigate the performance impact of parallel computing and vectorization on Matrix Multiplication ( $C = A \times B$ ). Matrix multiplication is a computationally intensive operation ( $O(N^3)$ ), making it an ideal candidate for high-performance optimization techniques.

We compare three specific implementations:

1. **Basic**: A standard naive implementation using three nested loops.
2. **Parallel**: An implementation using **OpenMP** to parallelize the outer loop across CPU cores.
3. **Vectorized (Optimized)**: An advanced implementation that combines:
  - **Multi-threading** via OpenMP.
  - **SIMD Intrinsics** (AVX2) to process multiple data points per instruction.
  - **Matrix Transposition** to optimize memory access patterns.

## 2 Theoretical Background

To understand the results presented in this report, it is necessary to define the optimization techniques used.

### 2.1 OpenMP and Multi-threading

OpenMP is an API that supports multi-platform shared-memory multiprocessing. It allows us to distribute iterations of a loop across multiple threads. In theory, using  $T$  threads should provide a speedup close to  $T \times$ . However, this is often limited by memory bandwidth and the overhead of creating threads.

### 2.2 SIMD (Single Instruction, Multiple Data)

Modern CPUs, such as the Intel i7-9750H used in this benchmark, support AVX2 (Advanced Vector Extensions). Standard code processes one number at a time (Scalar). SIMD allows the CPU to load 4 double precision numbers (256 bits) into a single register and multiply all of them simultaneously.

### 2.3 Cache Locality and Transposition

Matrices in C++ are stored in row-major order.

- Accessing  $A[i][k]$  is fast because we move linearly through memory.
- Accessing  $B[k][j]$  in the inner loop is slow because we jump  $N$  memory addresses for every step. This causes **Cache Misses**, where the CPU must wait for data to be fetched from RAM.

By transposing Matrix  $B$  ( $B^T$ ), we convert column access into row access, allowing the CPU to read contiguous memory blocks efficiently.

### 3 Methodology

The benchmarks were executed on an **Intel Core i7-9750H** (6 cores, 12 threads) running Arch Linux.

- **Compiler Flags:** `-O3 -march=native -fopenmp -mfma` (enables Fused Multiply-Add instructions).
- **Tools:** Linux `perf` was used to measure hardware counters, specifically **L3 Cache Misses** and **IPC** (Instructions Per Cycle).

### 4 Quantitative Results

This section presents the raw data collected during the benchmarks.

#### 4.1 Execution Time

Table 1 shows the execution time in seconds. While the Basic and Parallel versions grow exponentially with  $N$ , the Vectorized version remains highly efficient.

Size (N)	Basic (s)	Parallel (s)	Vectorized (s)
128	0.0037	0.0076	<b>0.0024</b>
256	0.0268	0.0082	<b>0.0041</b>
512	0.1962	0.0404	<b>0.0142</b>
1024	2.4993	0.3238	<b>0.0715</b>
2048	50.0731	19.8867	<b>0.5482</b>


*Table 1: Execution Time Comparison (seconds)*

#### 4.2 Computational Throughput (GFLOPS)

Table 2 measures the raw computing power utilized. The Vectorized approach consistently utilizes more of the CPU's potential.

Size (N)	Basic	Parallel	Vectorized
128	1.13	0.55	<b>1.73</b>
512	1.37	6.64	<b>18.90</b>
1024	0.86	6.63	<b>30.02</b>
2048	0.34	0.86	<b>31.34</b>

*Table 2: Throughput Comparison (GFLOPS)*




**Figure 1:** Computational Throughput (Higher is Better)

## 5 Performance Analysis

### 5.1 Execution Time and Overhead

As illustrated in Figure 2 and supported by Table 1, the performance gap widens significantly as the matrix size ( $N$ ) increases.


For small matrices ( $N = 128$ ), the **Parallel** implementation ( $0.0076s$ ) is actually slower than the **Basic** implementation ( $0.0037s$ ). This illustrates the *overhead of parallelism*: the time required to spawn OpenMP threads exceeds the time saved by dividing the small workload. However, at  $N = 2048$ , the Vectorized approach is roughly **91 times faster** than Basic.



**Figure 2:** Execution Time vs Matrix Size (Log Scale)

## 5.2 Analysis of the "Parallel Trap"

Ideally, adding more cores should improve performance linearly. However, Figure 3 shows that the Naive Parallel version hits a wall, achieving only a  $2.5 \times$  speedup at  $N = 2048$ .



**Figure 3:** Speedup Factor Relative to Basic Implementation


The reason for this poor scaling is found in Table 3. The Parallel version suffers from extreme *Cache Thrashing*. Because threads are reading Matrix B in columns, they load full cache lines but only use a single value before discarding them. At  $N = 2048$ , this

results in nearly **7 Billion cache misses**, forcing the CPU to stall while waiting for RAM.

**Table 3:** L3 Cache Misses at varying sizes (Lower is Better)

Size (N)	Basic	Parallel	Vectorized
512	283,197	320,573	<b>307,882</b>
1024	224,839,858	1,891,444	<b>2,018,394</b>
2048	7,927,259,288	6,912,059,223	<b>140,439,039</b>

As shown in Figure 4, the Vectorized version (which transposes Matrix B) reduces cache misses by a factor of 49 compared to the Parallel version (140M vs 6.9B).



**Figure 4:** L3 Cache Misses (Lower is Better)

## 6 Conclusion

This study demonstrates that high performance in Big Data applications requires a holistic approach to optimization.

- **Parallelization** alone provided a modest  $2.5\times$  gain but was severely bottlenecked by memory latency.
- **Vectorization + Memory Optimization** provided a massive  $91\times$  gain.

By aligning the data in memory (Transposition) to match the hardware's access patterns (SIMD/Cache Lines), we achieved a throughput of over **31 GFLOPS**, compared to less than 1 GFLOPS for the naive parallel approach.